

# Deep Residual Learning for Image Recognition

**Kaiming He - CVPR(2016)**

# CONTENTS

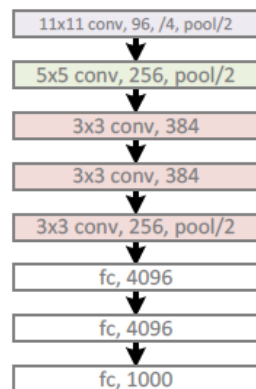
PART 01 - Problem

PART 02 - Network Architectures

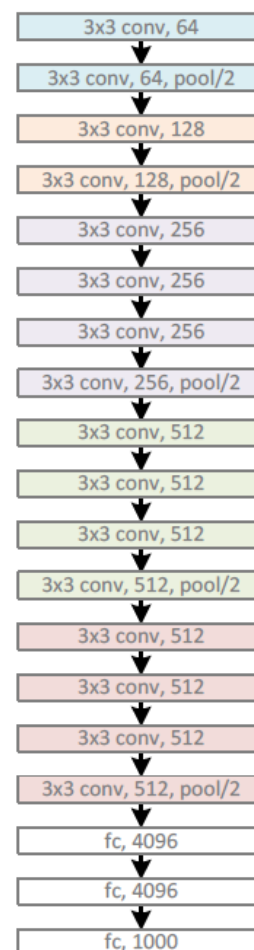
PART 03 - Experiments

# Revolution of Depth

AlexNet, 8 layers  
(ILSVRC 2012)



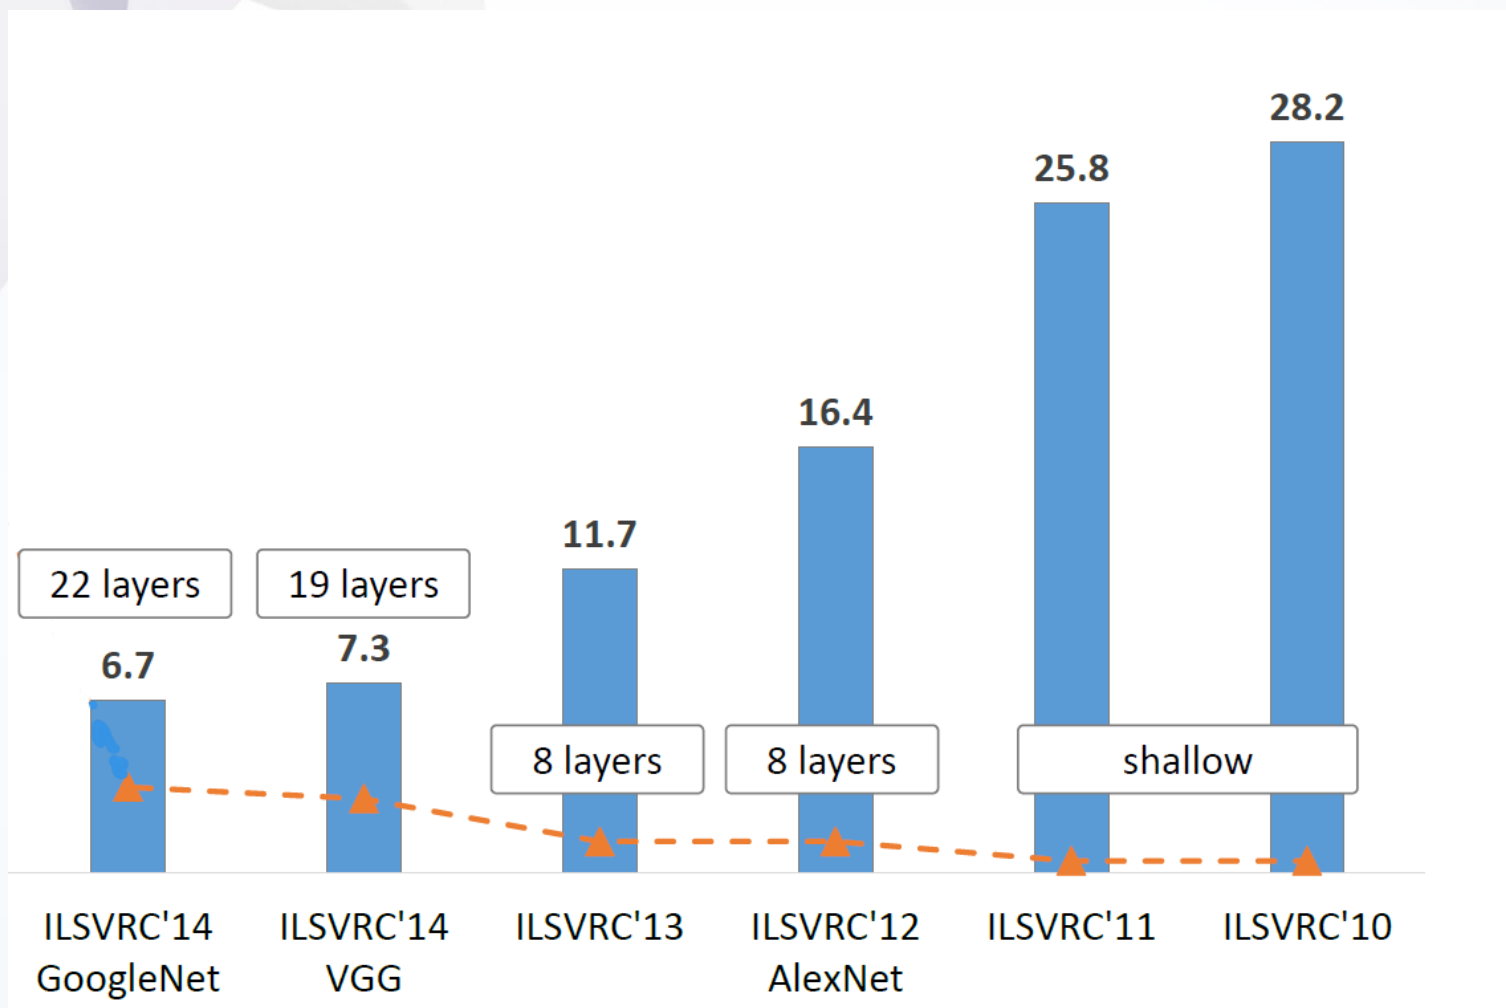
VGG, 19 layers  
(ILSVRC 2014)

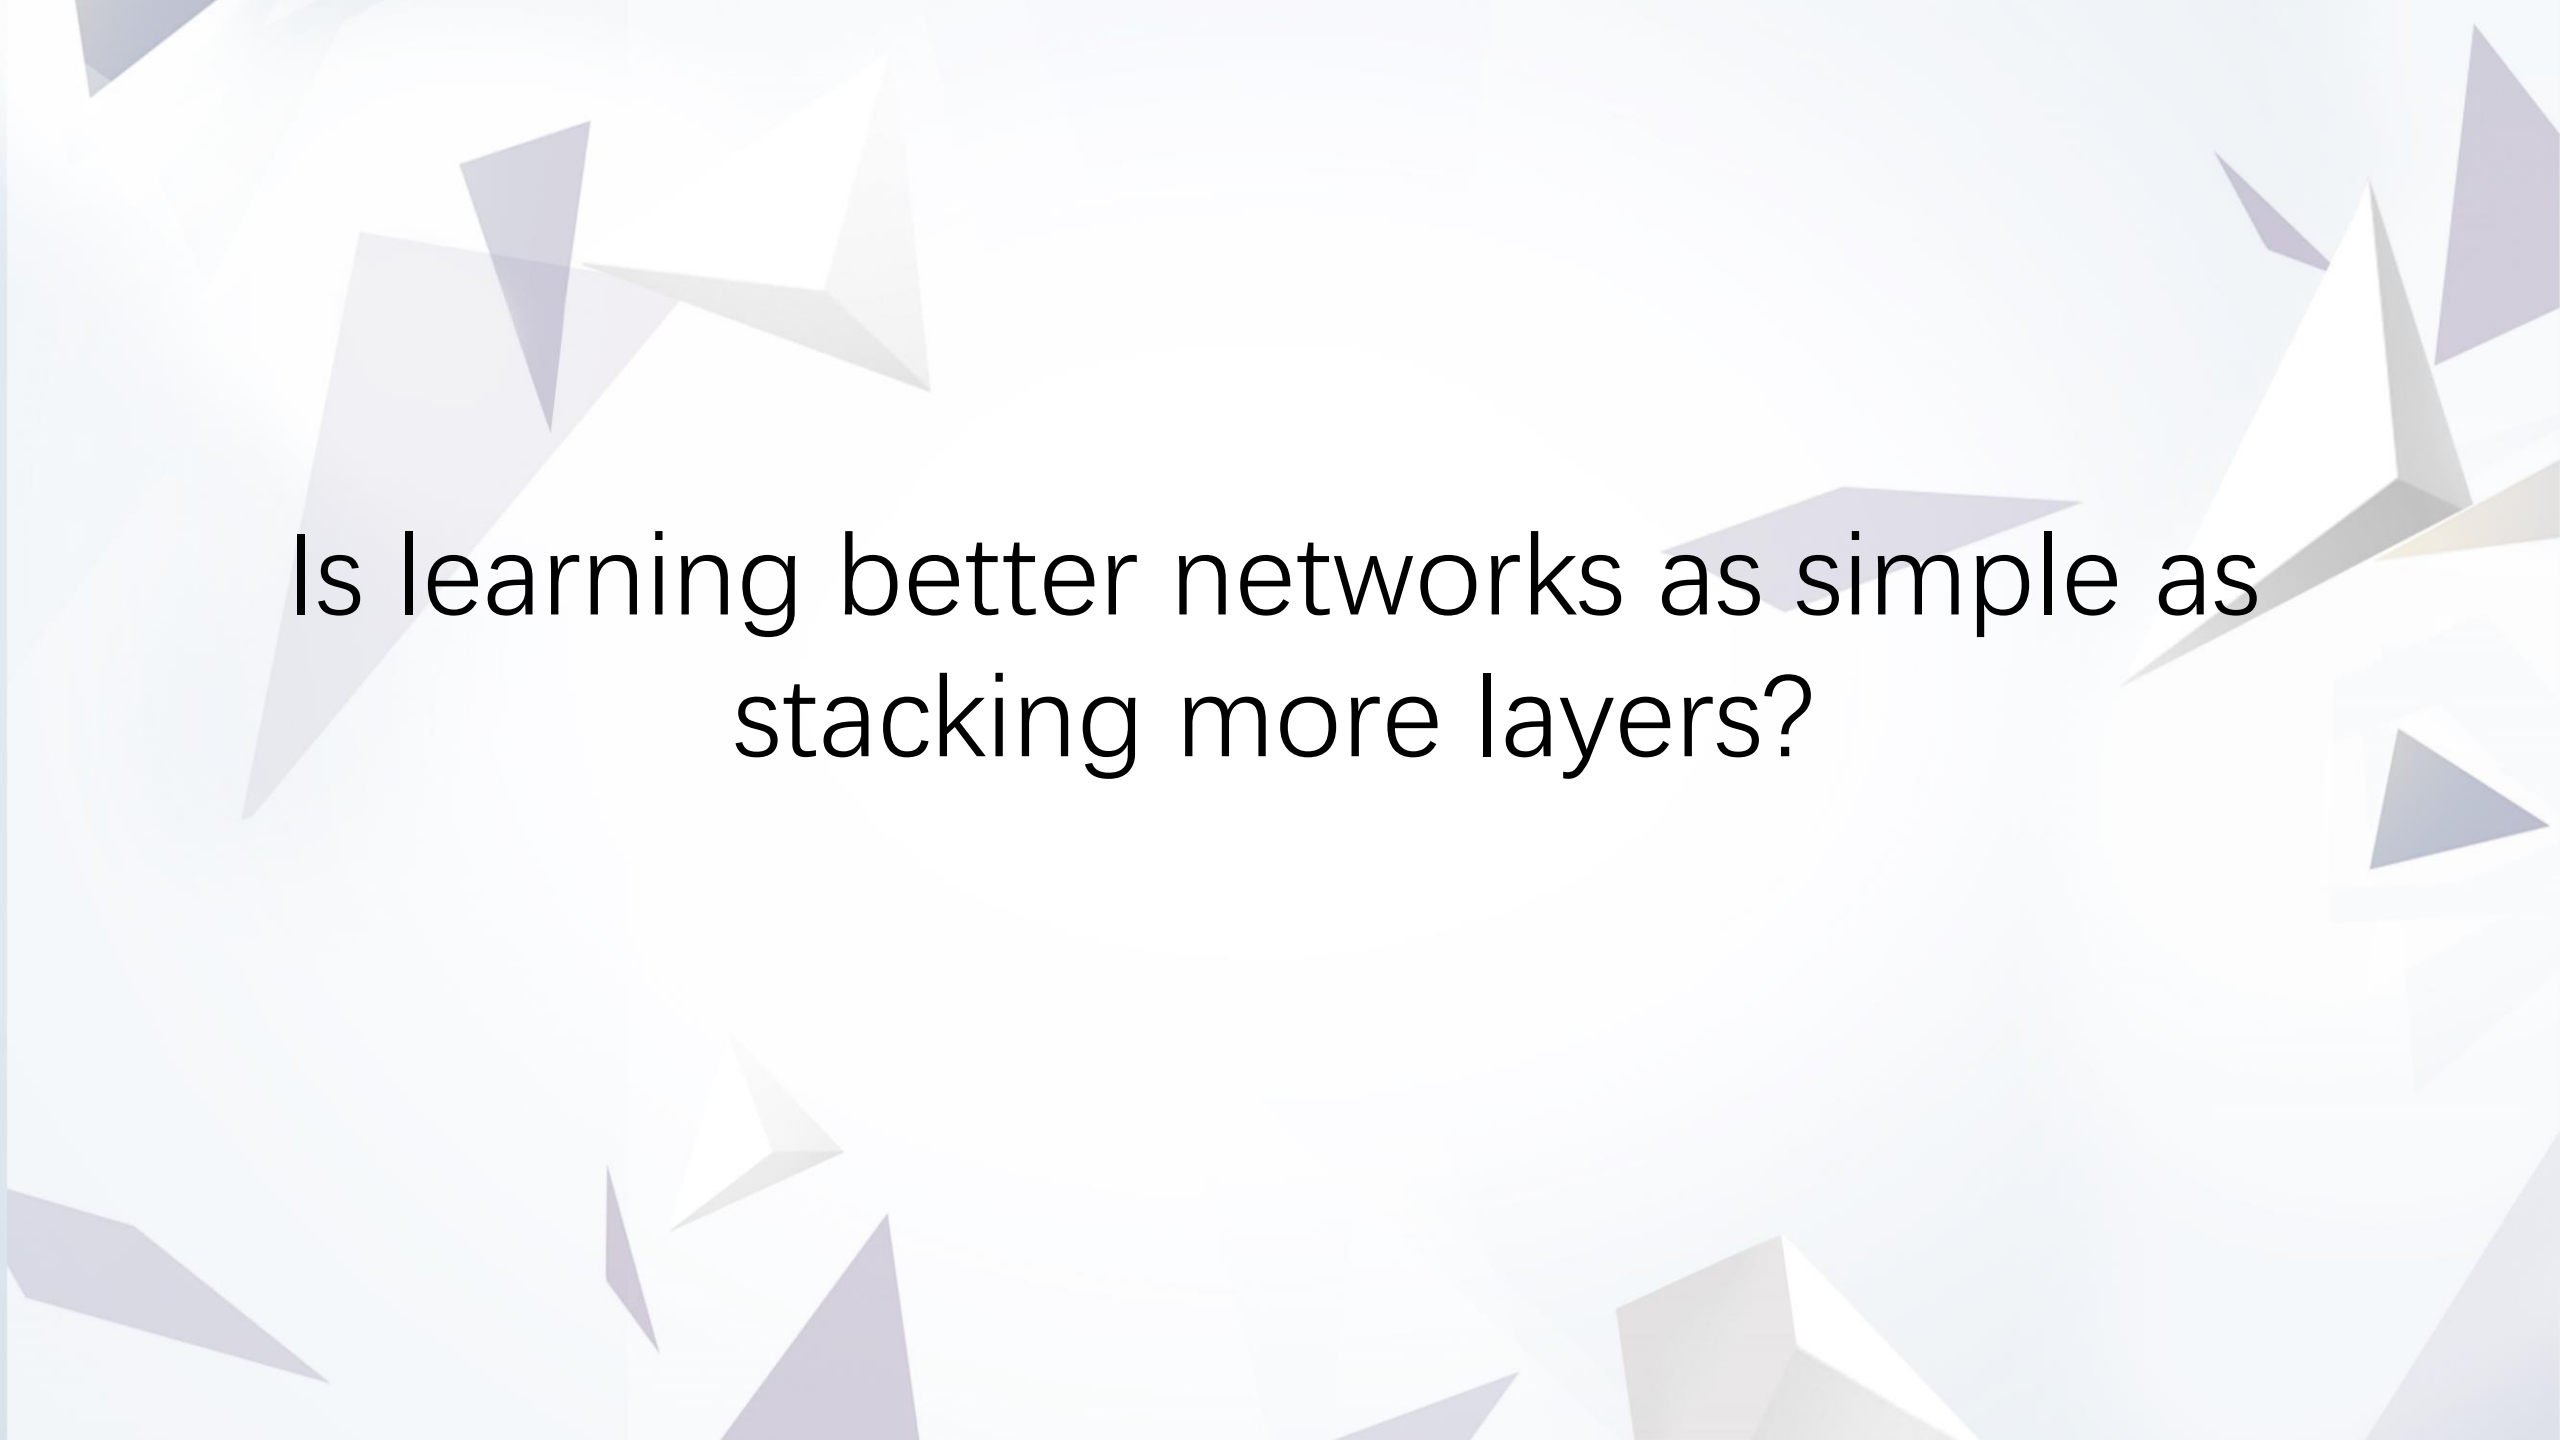


GoogleNet, 22 layers  
(ILSVRC 2014)



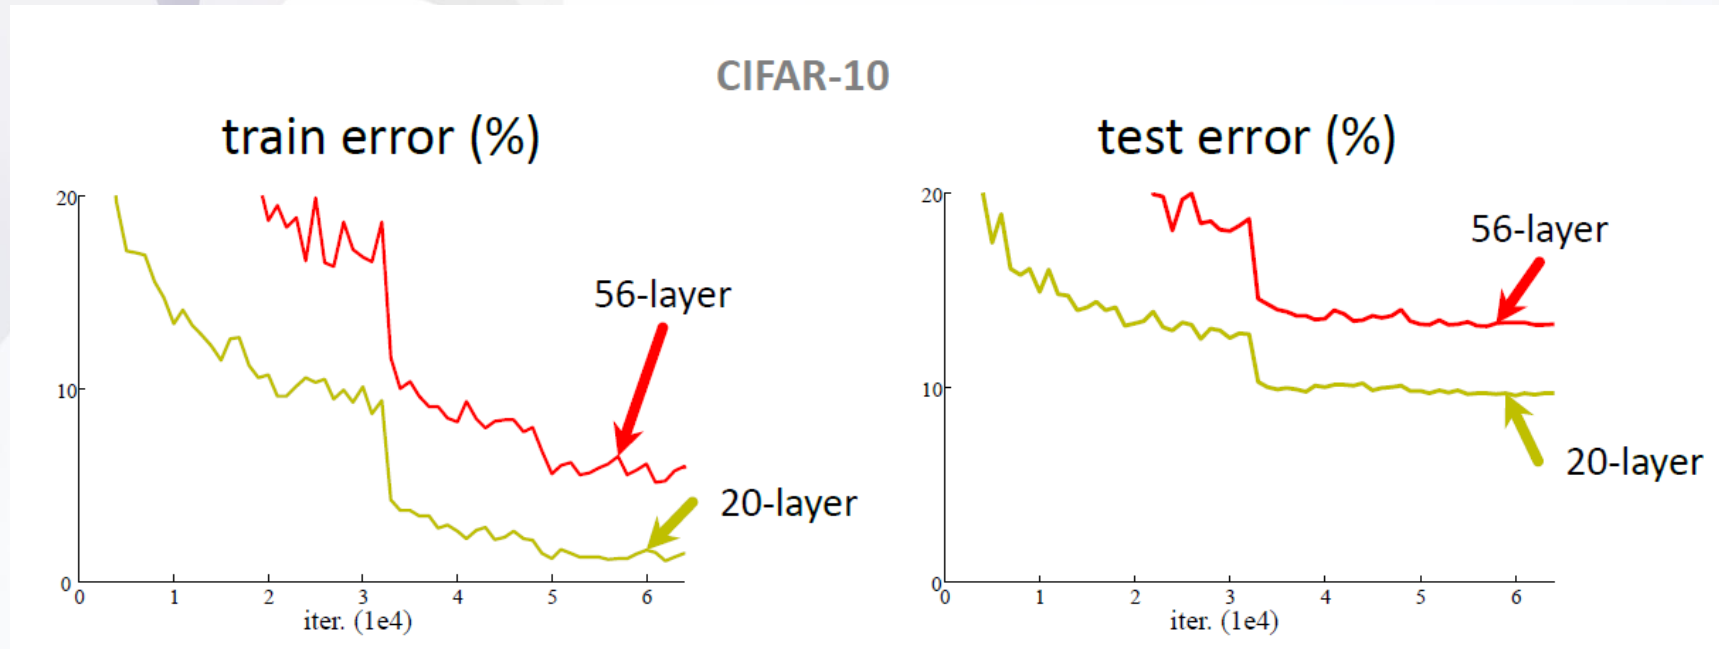
# Revolution of Depth





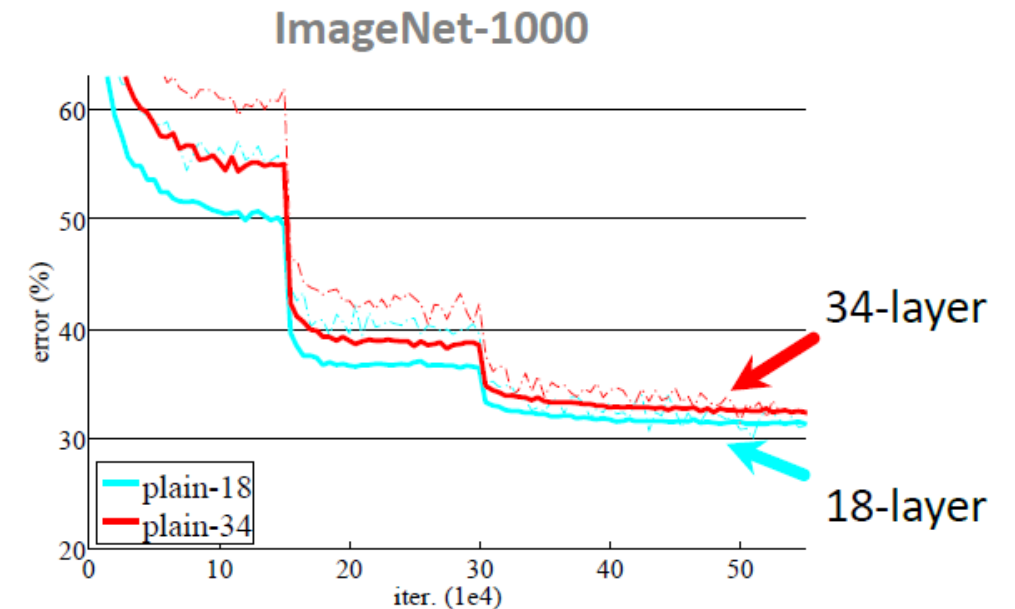
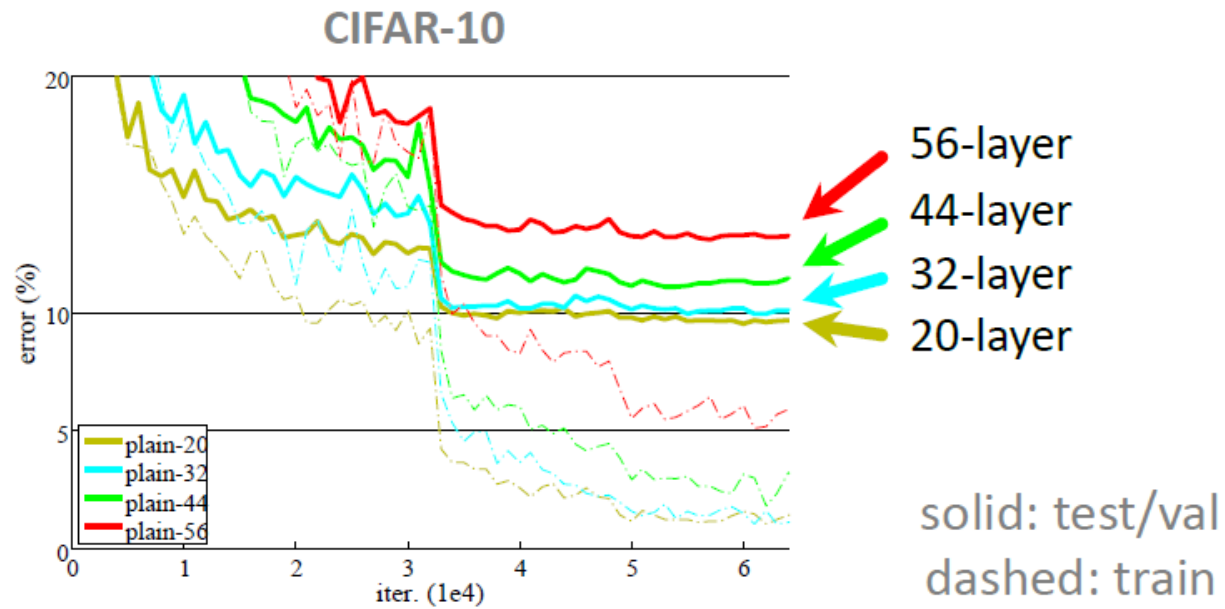
Is learning better networks as simple as  
stacking more layers?

# Simply stacking layers?



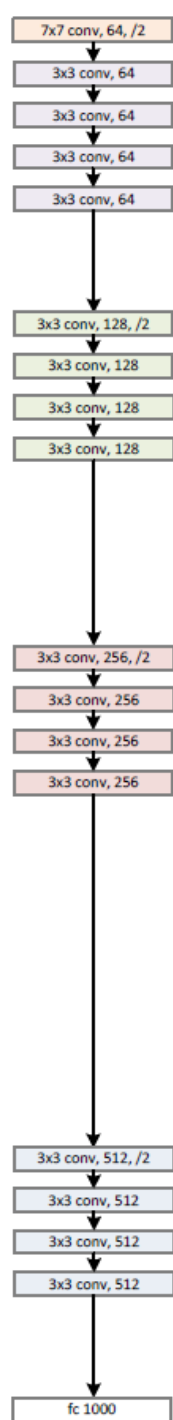
- plainnets: stacking 3x3 conv layers...
- 56-layer net has **higher training error** and test error than 20-layer net

# Simply stacking layers?

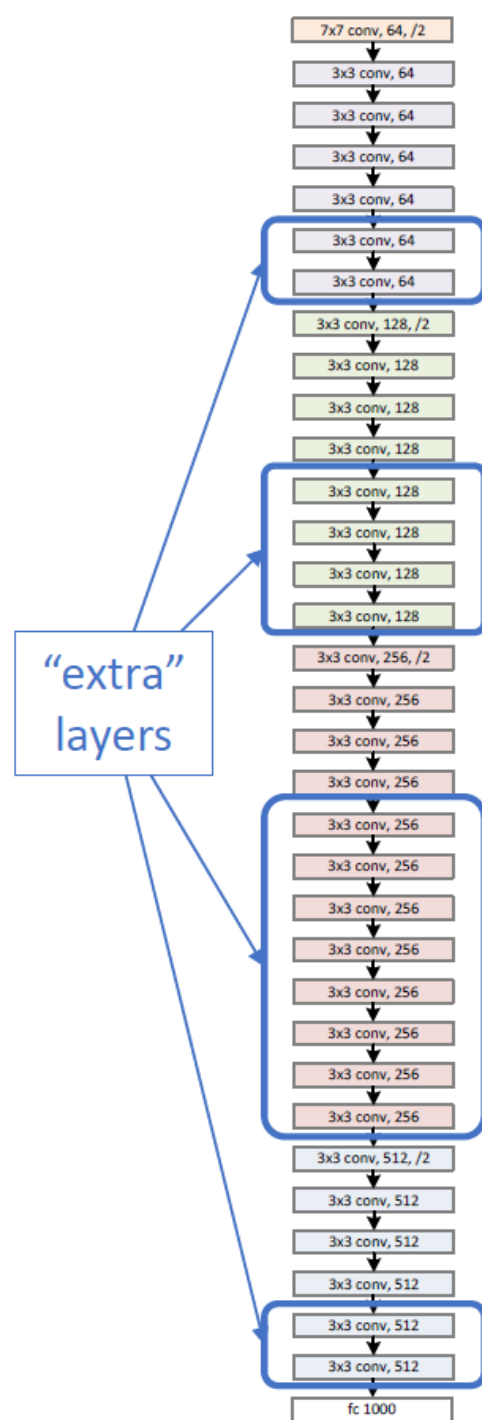


- “Overly deep” plain nets have **higher training error**
- A general phenomenon, observed in many datasets

a shallower  
model  
(18 layers)



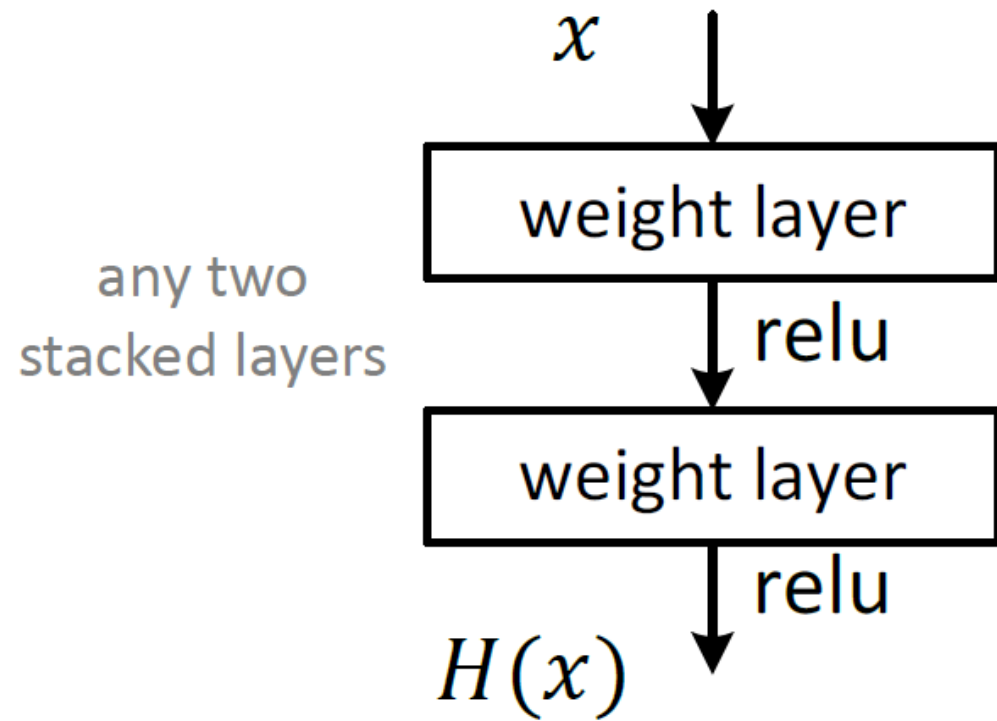
a deeper  
counterpart  
(34 layers)



- A deeper model should not have **higher training error**
- A solution by construction:
  - original layers: copied from a learned shallower model
  - extra layers: set as **identity**
  - at least the same training error
- **Optimization difficulties**: solvers cannot find the solution when going deeper...



# Plaint net



$H(x)$  is any desired mapping,  
hope the 2 weight layers fit  $H(x)$

# What is the residual?

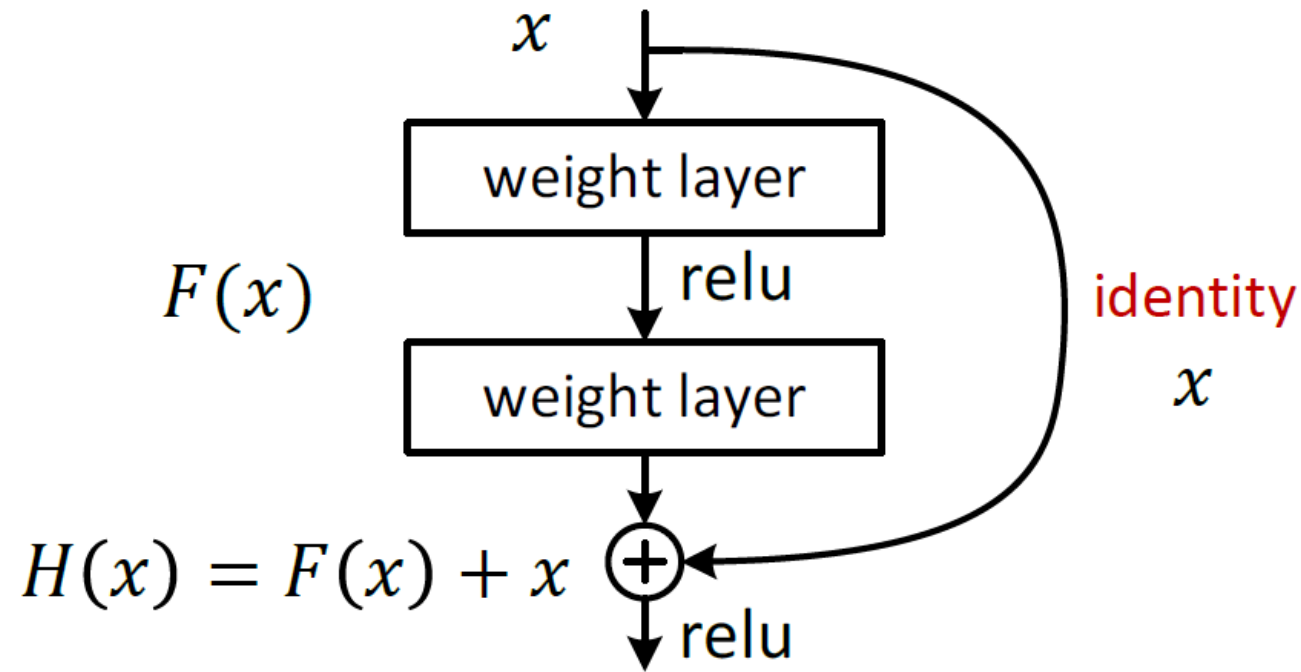
## **Definition(statistics):**

The difference between results obtained by observation and by computation from a formula or between the mean of several observations and any one of them

## **Example:**

Suppose we want to find an  $x$  such that  $f(x)=b$ , given an estimate  $x_0$  of  $x$ , the residual is  $b-f(x_0)$

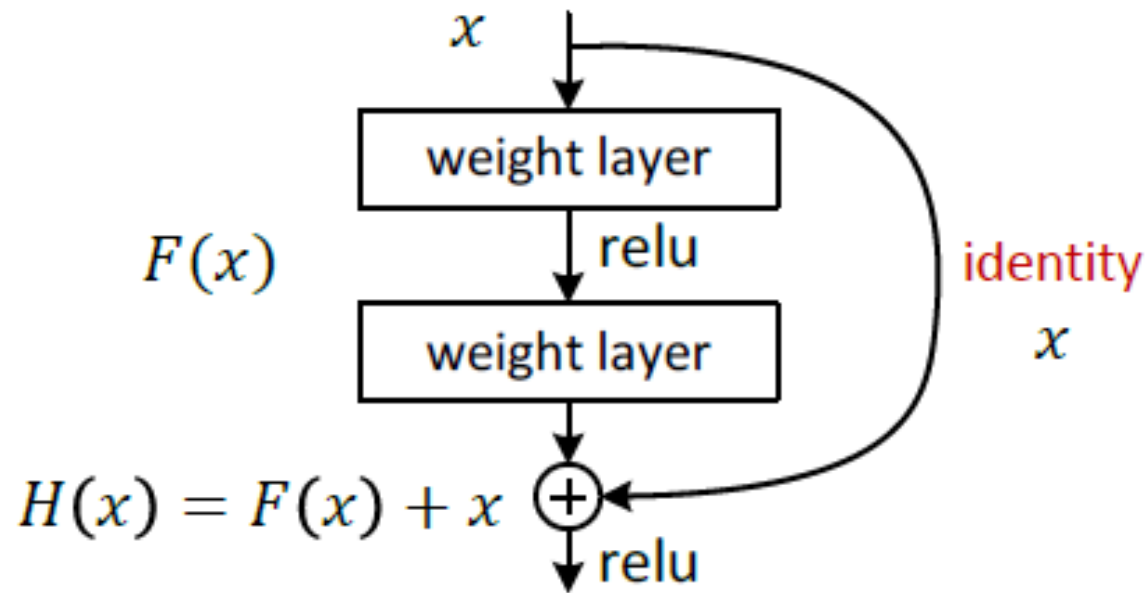
# Residual net



$H(x)$  is any desired mapping,  
~~hope the 2 weight layers fit  $H(x)$~~   
hope the 2 weight layers fit  $F(x)$   
let  $H(x) = F(x) + x$

# Residual net

- $F(x)$  is a **residual** mapping w.r.t. **identity**

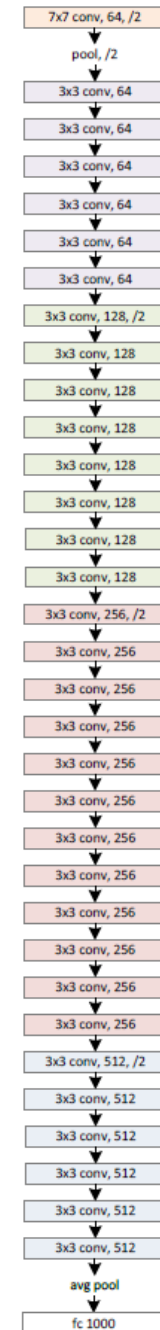


- If identity were optimal, easy to set weights as 0
- If optimal mapping is closer to identity, easier to find small fluctuations

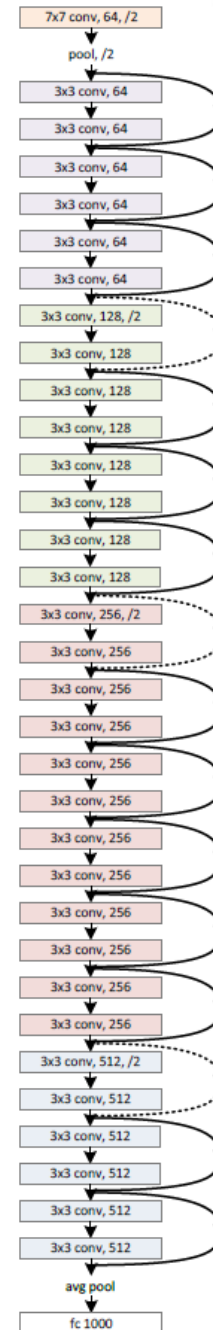
# Network “Design”

- Keep it simple
- Our basic design(VGG-style)
  - all 3x3 conv (almost)
  - spatial size /2 => # filters x2
  - Simple design; just deep!
- Other remarks:no max pooling (almost)
  - no hidden fc
  - no dropout

plain net



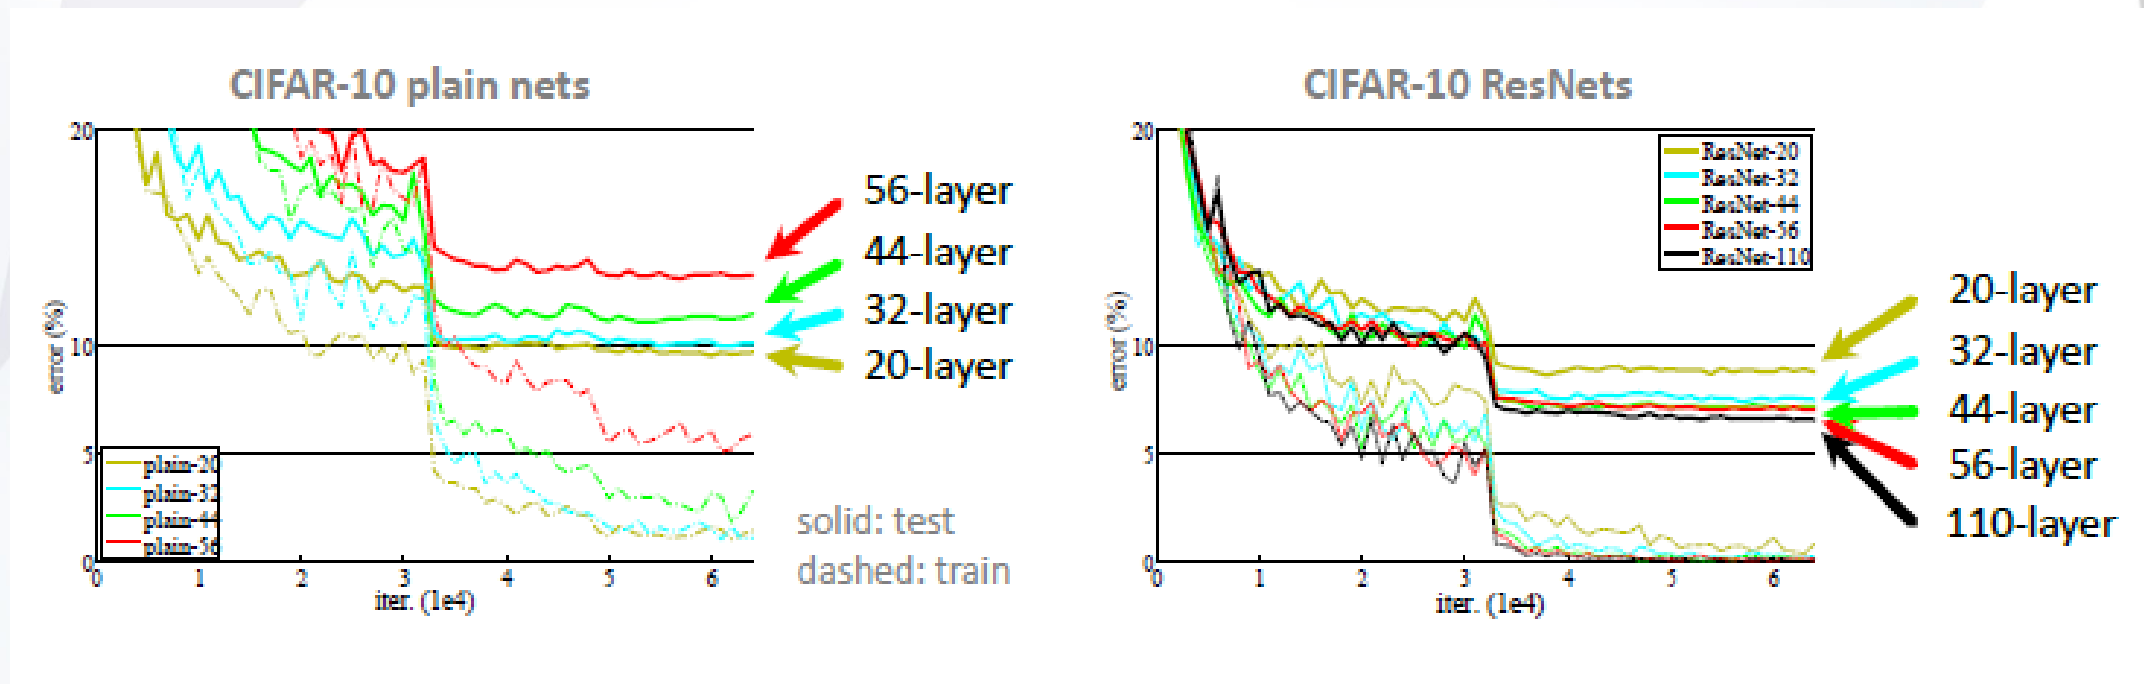
ResNet



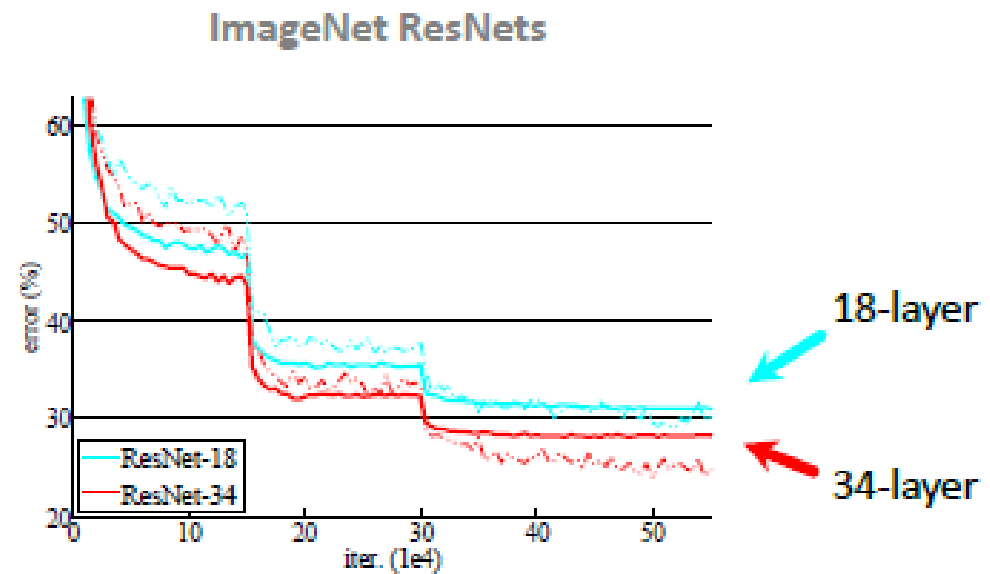
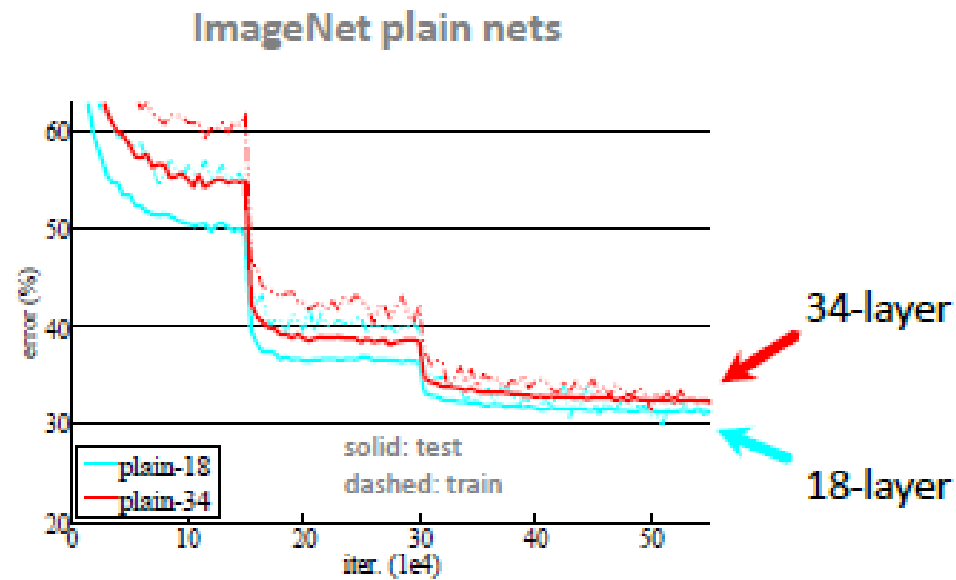
# Training

- All plain/residual nets are trained from scratch
- All plain/residual nets use Batch Normalization
- Standard hyper-parameters & augmentation

# CIFAR-10 experiments



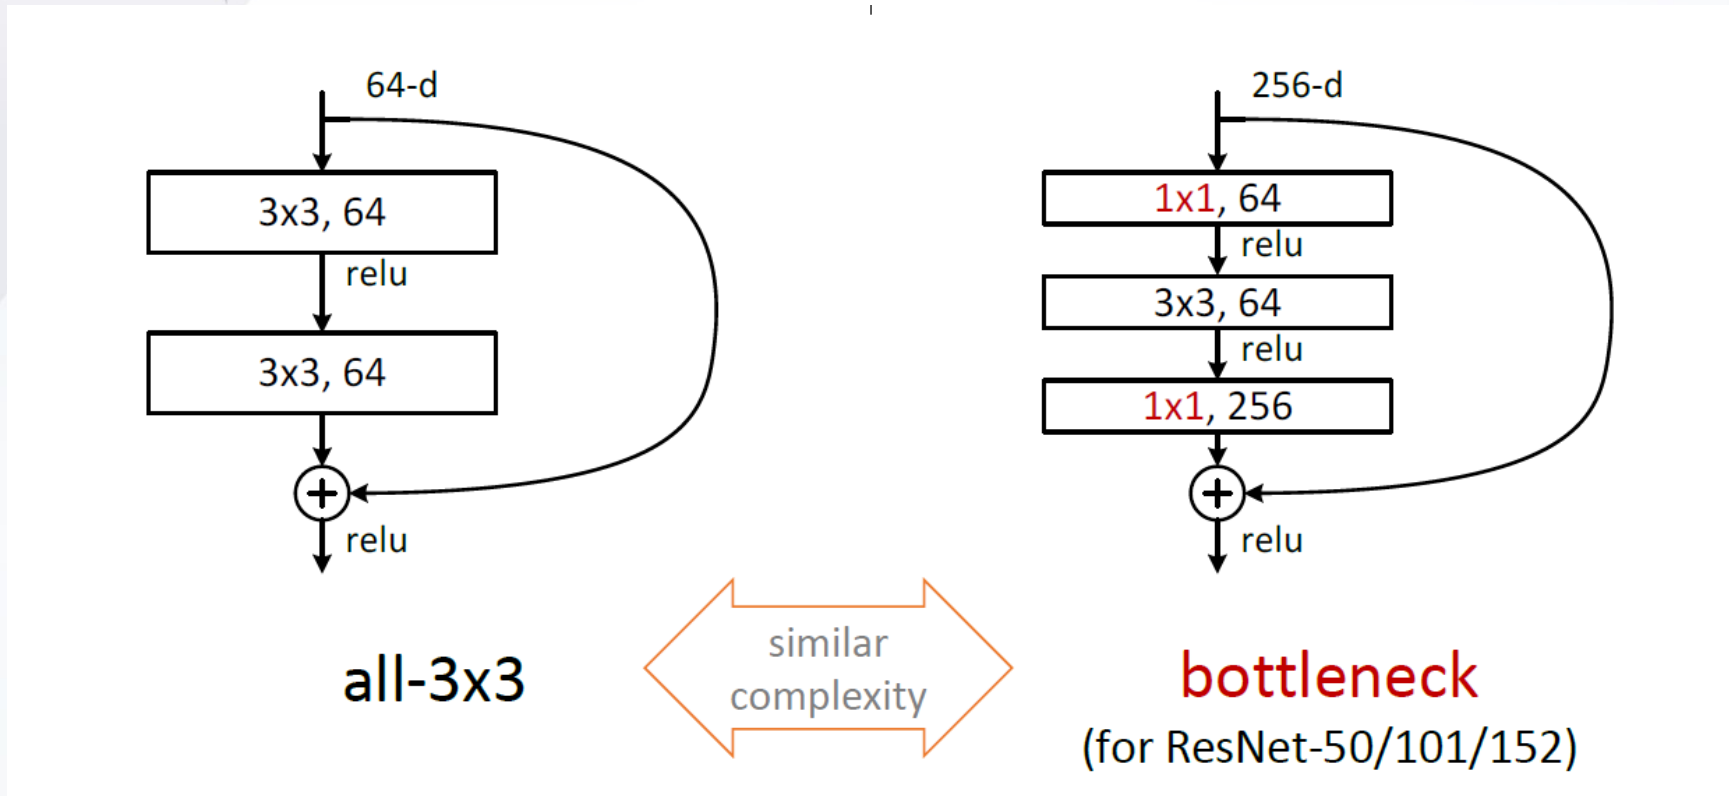
# ImageNet experiments



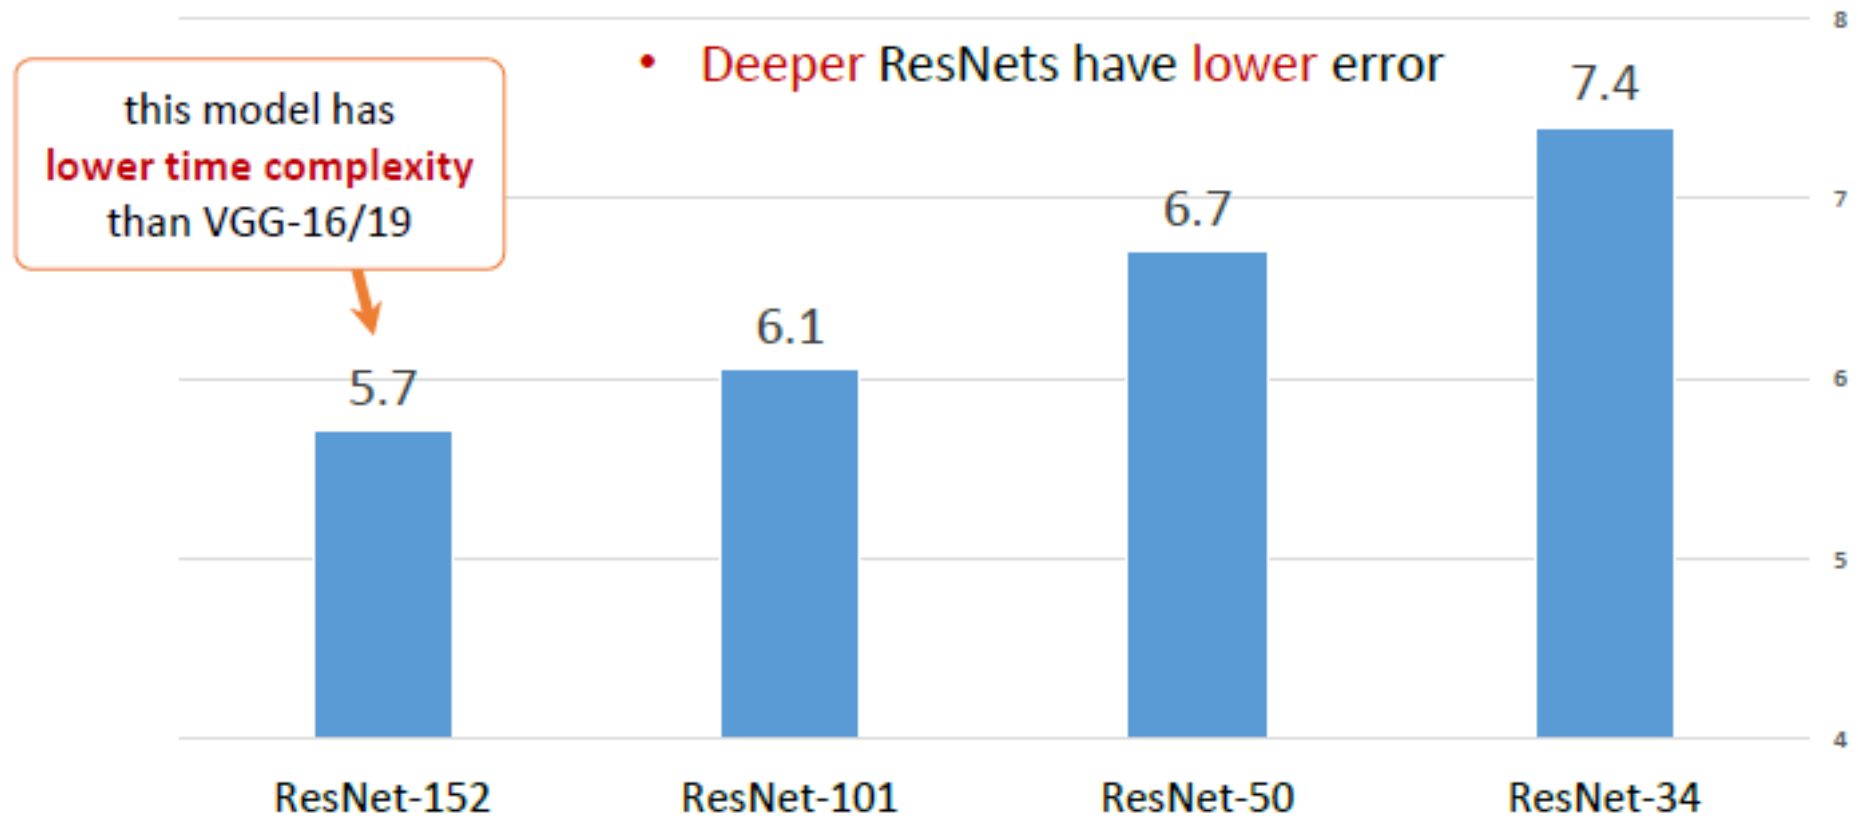
- Deep ResNets can be trained without difficulties
- Deeper ResNets have **lower training error**, and also lower test error



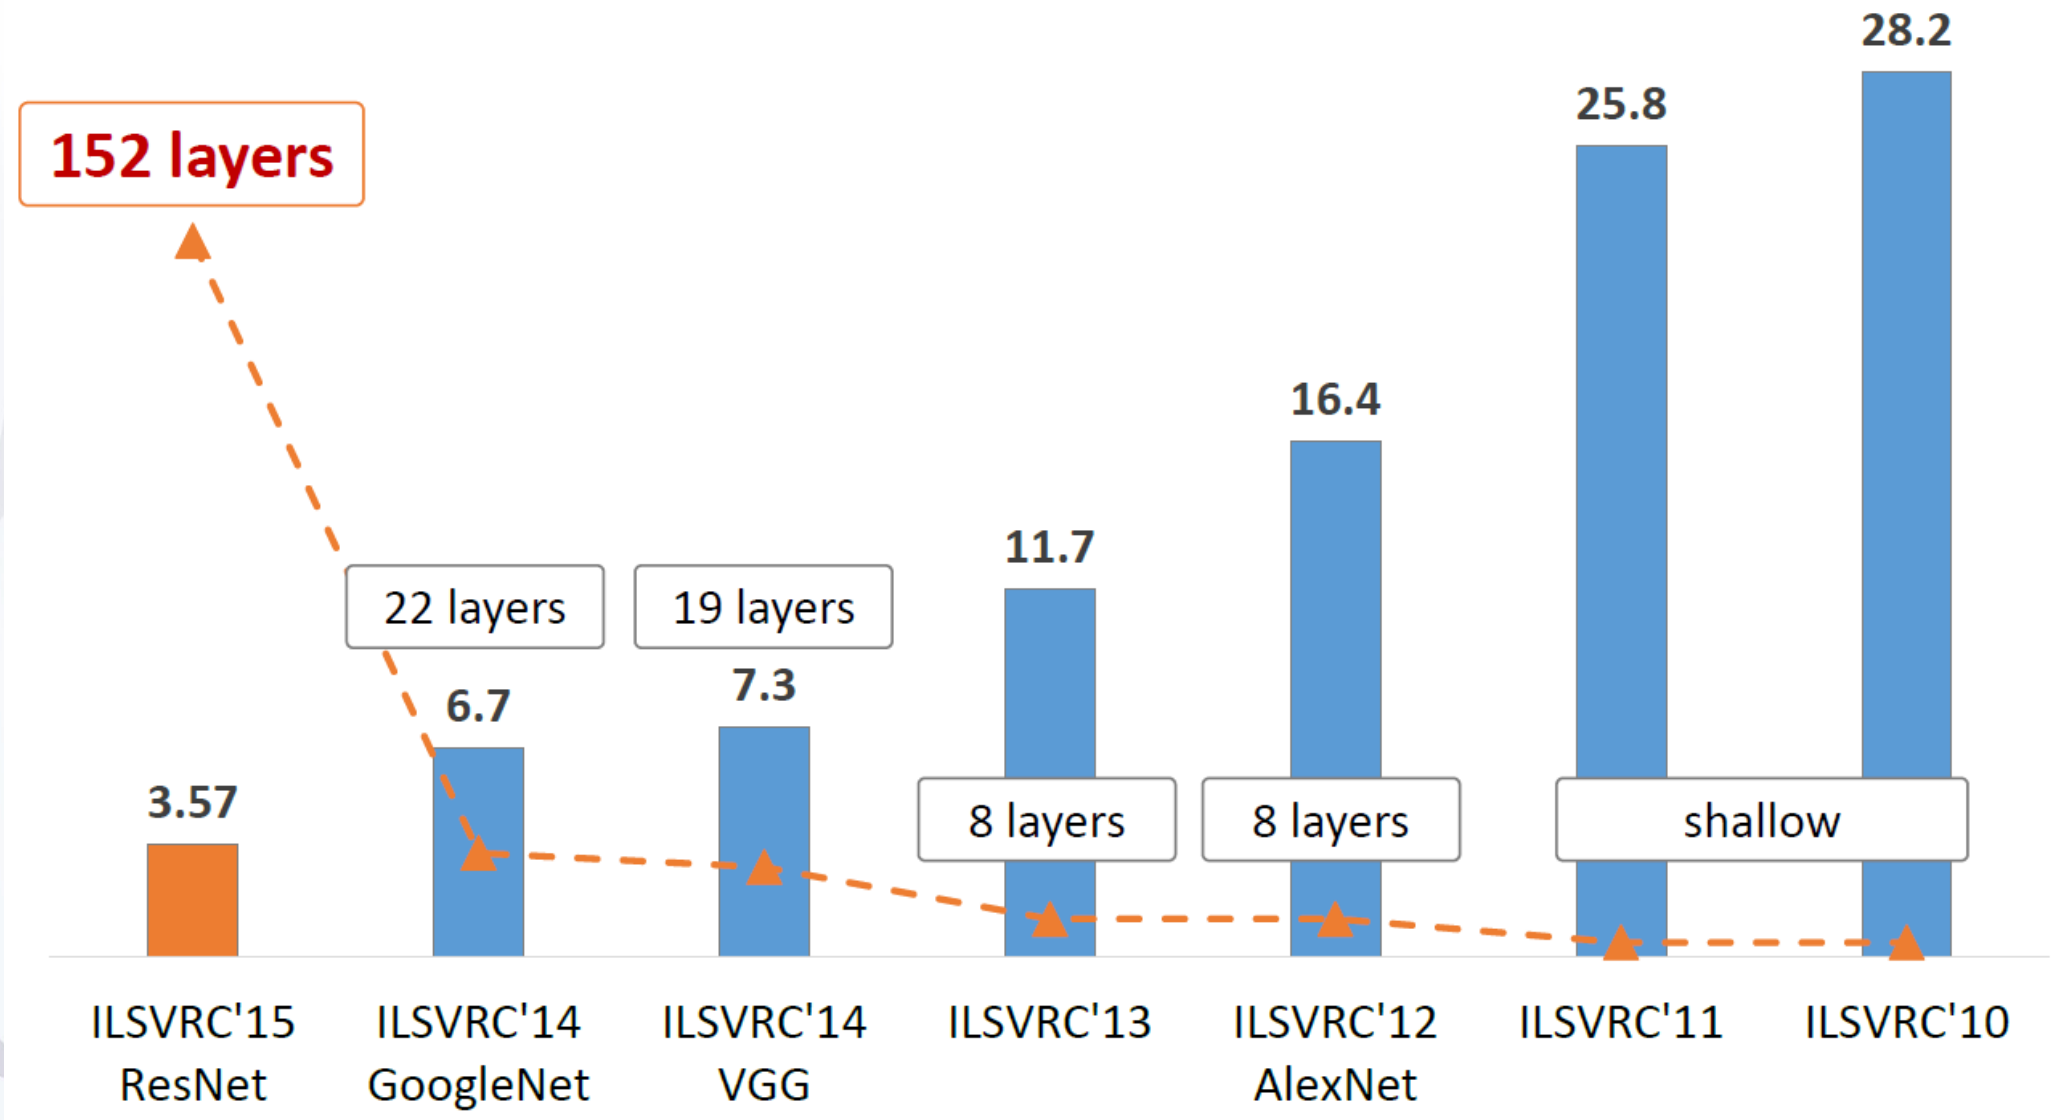
# A practical design of going deeper



# ImageNet experiments



# ImageNet experiments



# Why residual works?

The main reason the residual network works is that it's so easy for these extra layers to learn the identity function that you're kind of guaranteed that it doesn't hurt performance. And then lot of time you maybe get lucky and even helps performance, or at least is easier to go from a decent baseline of not hurting performance, and then creating the same can only improve the solution from there.

$$z^{[l+1]} = W^{[l+1]}a^{[l]} + b^{[l+1]}$$

$$a^{[l+1]} = g(z^{[l+1]})$$

$$z^{[l+2]} = W^{[l+2]}a^{[l+1]} + b^{[l+2]}$$

$$a^{[l+2]} = g(z^{[l+2]}) \quad \longrightarrow \quad a^{[l+2]} = g(W^{[l+2]}a^{[l+1]} + b^{[l+2]} + a^{[l]})$$

# Personal experiment

```
Test - > Loss: 0.414 | Acc: 86.237% (6554/7600)
Test - > Loss: 0.415 | Acc: 86.260% (6642/7700)
Test - > Loss: 0.417 | Acc: 86.154% (6720/7800)
Test - > Loss: 0.416 | Acc: 86.152% (6806/7900)
Test - > Loss: 0.418 | Acc: 86.112% (6889/8000)
Test - > Loss: 0.416 | Acc: 86.148% (6978/8100)
Test - > Loss: 0.417 | Acc: 86.146% (7064/8200)
Test - > Loss: 0.418 | Acc: 86.145% (7150/8300)
Test - > Loss: 0.417 | Acc: 86.179% (7239/8400)
Test - > Loss: 0.418 | Acc: 86.129% (7321/8500)
Test - > Loss: 0.419 | Acc: 86.093% (7404/8600)
Test - > Loss: 0.419 | Acc: 86.023% (7484/8700)
Test - > Loss: 0.420 | Acc: 86.000% (7568/8800)
Test - > Loss: 0.419 | Acc: 85.989% (7653/8900)
Test - > Loss: 0.419 | Acc: 85.956% (7736/9000)
Test - > Loss: 0.418 | Acc: 85.967% (7823/9100)
Test - > Loss: 0.416 | Acc: 86.022% (7914/9200)
Test - > Loss: 0.416 | Acc: 86.032% (8001/9300)
Test - > Loss: 0.415 | Acc: 85.979% (8082/9400)
Test - > Loss: 0.414 | Acc: 85.979% (8168/9500)
Test - > Loss: 0.414 | Acc: 85.990% (8255/9600)
Test - > Loss: 0.414 | Acc: 85.979% (8340/9700)
Test - > Loss: 0.415 | Acc: 85.959% (8424/9800)
Test - > Loss: 0.416 | Acc: 85.949% (8509/9900)
[=====]
Test - > Loss: 0.416 | Acc: 85.940% (8594/10000)
```

Epoches:5000

ResNet-18

Simple preprocess

# Further study

<http://kaiminghe.com/icml16tutorial/index.html>

## **Publications:**

- [a] [Deep Residual Learning for Image Recognition](#)  
Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. CVPR 2016.
- [b] [Identity Mappings in Deep Residual Networks](#)  
Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Technical report, arXiv 2016.

## **Resources:**

- [tutorial slides](#)
- [slides](#) and [video](#) for the talk at ICCV 2015 ImageNet and COCO joint workshop.
- [code/models](#) of 50, 101, and 152-layer ResNets pre-trained on ImageNet.
- [code](#) of 1001-layer ResNet on CIFAR.
- [list](#) of third-party ResNet implementations on ImageNet, CIFAR, MNIST, etc.



Thanks